



It's a totally valid feature request though, especially given that R8/AppReduce can do optimizations that would invalidate measurements.

Do you have a preference as to which approach would be preferable?

IIRC, blackhole is typically passing to a noop native function, where the optimizer can't make assumptions, though could also be done with some sort of @Keep for R8/appreduce specific op

Caliper's return approach could theoretically be lower overhead for tiny microbenchmarks, which I like, but I wonder if it would be easy to accidentally misuse. Primarily, I'd worry about perfor



**ma...@google.com** <ma...@google.com> [#3](#)

My preference would be a BlackHole approach, seems easier to use correctly.

My only trouble with recommending a blackhole is I'm not totally sure how they work, or sure how much the CPU (even at the microcode, say) might be able to reason that the blackhole is dead code and skip some computations. But I don't even know if that's really a thing that happens. All I know is I should distrust every layer in the stack because they all seem to be full of sneaky optimisers and caches :-)

On Thu, Sep 16, 2021 at 9:02 AM ccraik <[buganizer-system+ccraik@google.com](mailto:buganizer-system+ccraik@google.com)> wrote:

[- Show quoted text -](#)



**cc...@google.com** <cc...@google.com> [#4](#)

My understanding is the straightforward implementation of a blackhole is to throw it at a noop JNI method. Just declare a `native void blackhole(Object arg)`, and an empty C method c  
Unfortunately the nature of gradle lacking build support for R8 with our recommended config, I'm not able to test something like that on our side.



**ma...@google.com** <ma...@google.com> [#5](#)

Is there much overhead to a JNI call? Just curious if you know, I don't.

On Wed, Oct 13, 2021 at 12:36 PM ccraik <[buganizer-system+ccraik@google.com](mailto:buganizer-system+ccraik@google.com)> wrote:

[- Show quoted text -](#)